
Chapter Three (2/2)

1

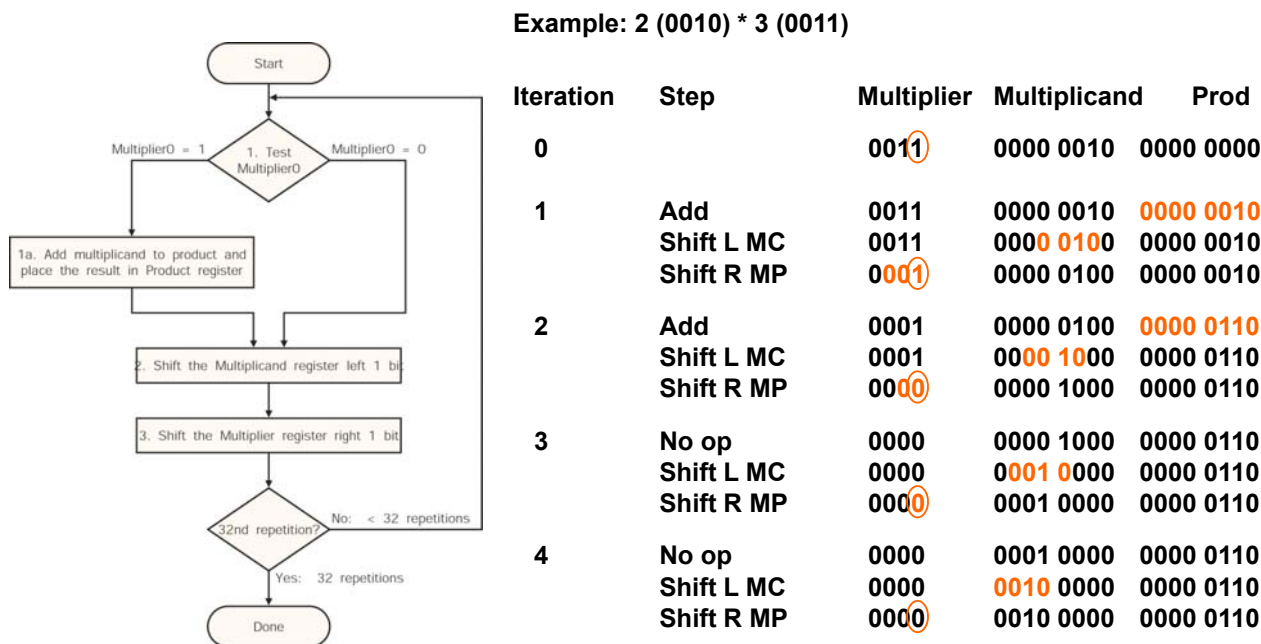
Multiplication

- More complicated than addition
 - accomplished via shifting and addition
- n bits \times n bits = $2n$ bit product
- Binary makes it easy:
 - 0 => place 0 (0 x multiplicand)
 - 1 => place a copy (1 x multiplicand)

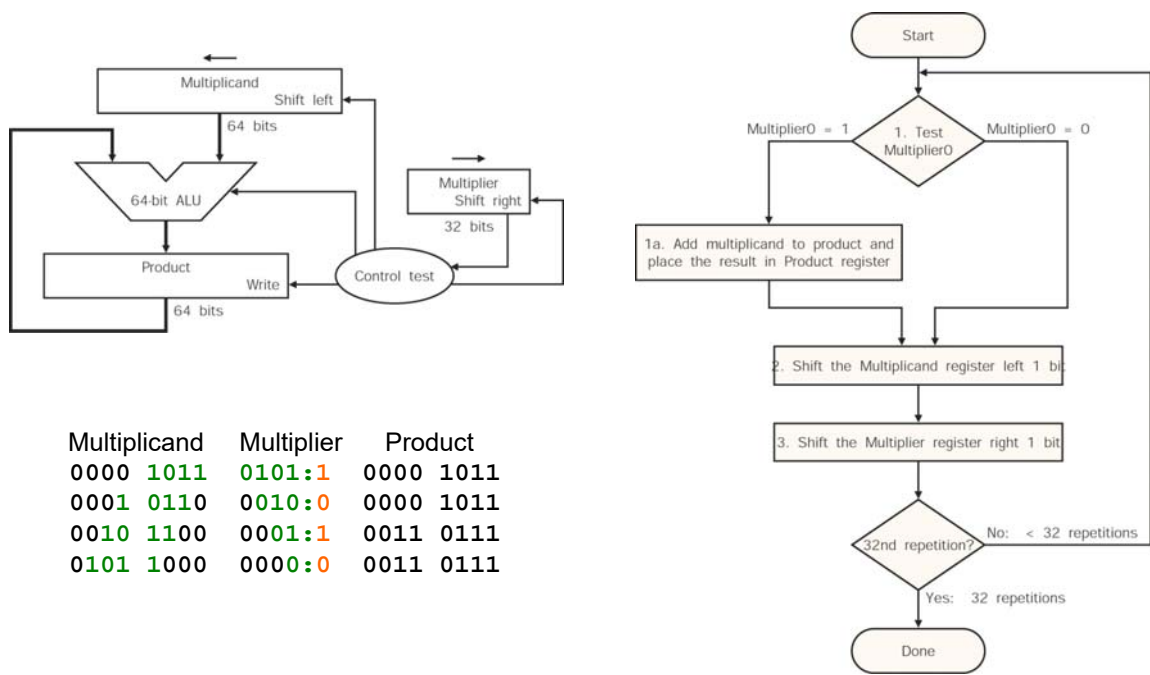
```
      0010 (multiplicand)
    x 1011 (multiplier)
    -----
      0010
     0010
    0000
   0010
  -----
  0010110
```

- Negative numbers: convert and multiply
 - there are better techniques, we won't look at them

Multiplication: Version 1

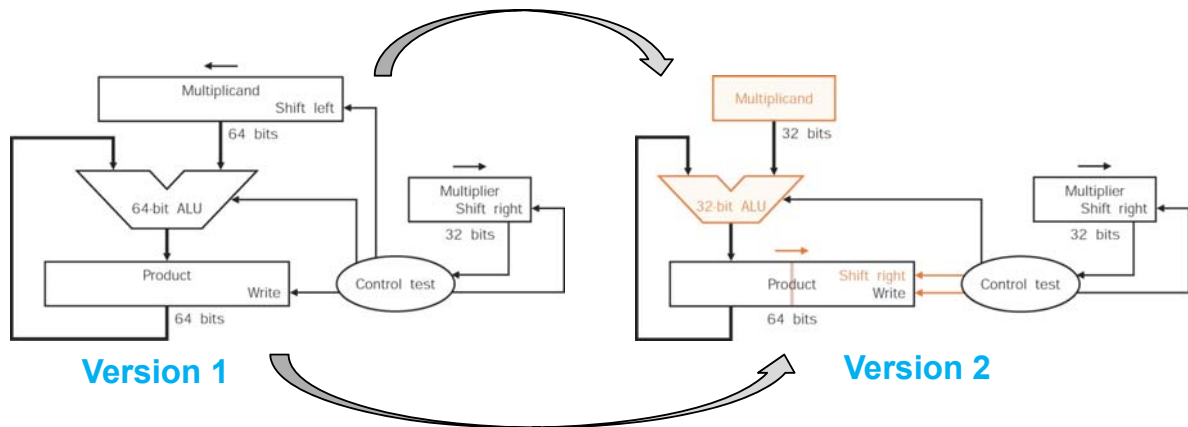


Multiplication: Version 1



Observations on Multiply Version 1

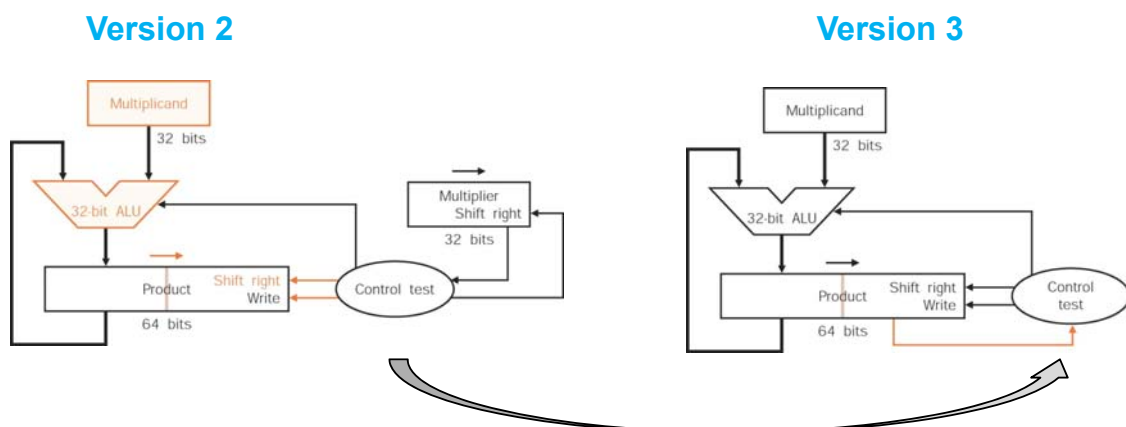
- 1 clock per cycle => 32 clocks per multiply
- 1/2 bits in multiplicand always 0
=> 64-bit adder is wasted
- 0's inserted in left of multiplicand as shifted
=> least significant bits of product never changed once formed
- **Instead of shifting multiplicand to left, shift product to right.**



CE, KWU Prof. S.W. LEE 5

Observations on Multiply Version 2

- **2 cycles per bit** because shift is performed after addition
- Product register wastes space that exactly matches size of multiplier
- Both Multiplier register and Product register require right shift
- Combine Multiplier register and Product register



CE, KWU Prof. S.W. LEE 6


Observations on Multiply Version 3

- **2 cycles per bit** because shift is performed after addition
- MIPS registers Hi and Lo are left and right half of Product
- Gives us MIPS instruction MultU
- What about signed multiplication?
 - easiest solution is to make both positive & remember whether to complement product when done (leave out the sign bit, run for 31 steps)
 - apply definition of 2's complement
 - need to sign-extend partial products and subtract at the end
 - Booth's Algorithm is elegant way to multiply signed numbers using same hardware as before and save cycles
 - can be modified to handle multiple bits at a time

CE, KWU Prof. S.W. LEE 7

Faster Multiplier

- Uses multiple adders
 - Cost/performance tradeoff



- Can be pipelined
 - Several multiplication performed in parallel


CE, KWU Prof. S.W. LEE 8

MIPS Multiplication

- Two 32-bit registers for product
 - HI: most-significant 32 bits
 - LO: least-significant 32-bits
- Instructions
 - **mult rs, rt / multu rs, rt**
 - 64-bit product in HI/LO
 - **mfhi rd / mflo rd**
 - Move from HI/LO to rd
 - Can test HI value to see if product overflows 32 bits
 - **mul rd, rs, rt**
 - Least-significant 32 bits of product → rd

Booth's Algorithm

- Multiplier



- $a \times 30 = a \times (32 - 2)$

- Depending on the current and previous bits, do one of the following:
 - 00: no operation
 - 01: add the multiplicand to the left half of the product
 - 10: subtract the multiplicand from the left half of the product
 - 11: no operation

Booth's Algorithm

- Depending on the current and previous bits, do one of the following:
 - 00: no operation
 - 01: add the multiplicand to the left half of the product
 - 10: subtract the multiplicand from the left of the product
 - 11: no operation

Example: 2 (0010) * 3 (0110)


Iteration	Step	Multiplicand	Prod
0		0010	0000 0110 0
1	00 -> no op Shift R MP	0010 0010	0000 0110 0 0000 0011 0
2	10 -> sub Shift R MP	0010 0010	1110 0011 0 1111 0001 1
3	11 -> no op Shift R MP	0010 0010	1111 0001 1 1111 1000 1
4	01 -> add Shift R MP	0010 0010	0001 1000 1 0000 1100 0

CE, KWU Prof. S.W. LEE 11

Radix-4 Modified Booth's Algorithm

- Considering three consecutive bits:

x_i	x_{i-1}	x_{i-2}	y_i	y_{i-1}	operation	comments
0	0	0	0	0	+0	string of zeros
0	1	0	0	1	+A	a single 1
1	0	0	1	0	-2A	beginning of 1's
1	1	0	0	1	-A	beginning of 1's
0	0	1	0	1	+A	end of 1's
0	1	1	1	0	+2A	end of 1's
1	0	1	0	1	-A	a single 0
1	1	1	0	0	+0	string of 1's



- Further improvements
 - Radix-8 modified Booth's algorithm
 - Other parallel approaches

CE, KWU Prof. S.W. LEE 12

Modified Booth's Algorithm

- Depending on the current and previous bits, do one of the following:

x_i	x_{i-1}	x_{i-2}	y_i	y_{i-1}	operation	comments
0	0	0	0	0	+0	string of zeros
0	1	0	0	1	+A	a single 1
1	0	0	$\bar{1}$	0	-2A	beginning of 1's
1	1	0	0	$\bar{1}$	-A	beginning of 1's
0	0	1	0	1	+A	end of 1's
0	1	1	1	0	+2A	end of 1's
1	0	1	0	$\bar{1}$	-A	a single 0
1	1	1	0	0	+0	string of 1's

Example: 2 (0010) * 3 (0110)

Iteration	Step	Multiplicand	Prod
0		0010	0000 0110 0
1	100 → -2A Shift R2 MP	0010 0010	1100 0110 0 1111 0001 1
2	011 → 2A Shift R2 MP	0010 0010	0011 0001 1 0000 1100 0

CE, KWU Prof. S.W. LEE 13

Unsigned Divide: Paper & Pencil

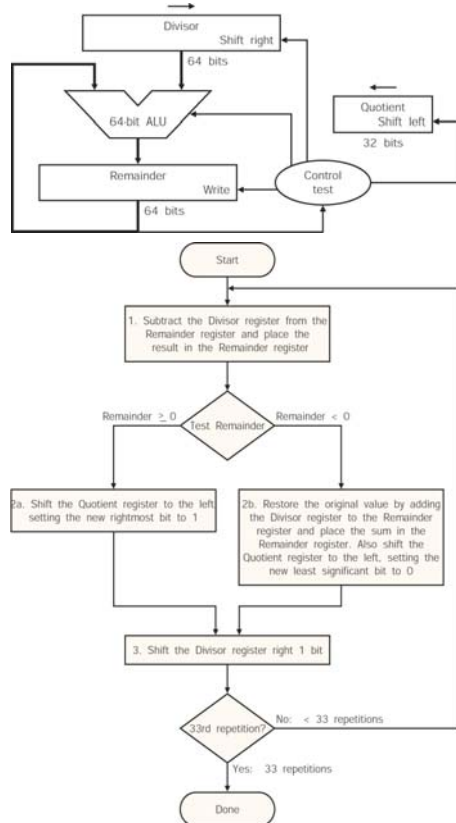
Divisor	1000	$\overline{1001}$	Quotient
		1001010	Dividend
		$\underline{-1000}$	
		10	
		101	
		$\underline{1010}$	
		$\underline{-1000}$	
		10	Remainder (or Modulo result)

- See how big a number can be subtracted, creating quotient bit on each step
Binary => 1 * divisor or 0 * divisor

Dividend = Quotient x Divisor + Remainder

CE, KWU Prof. S.W. LEE 14

Division: Version 1



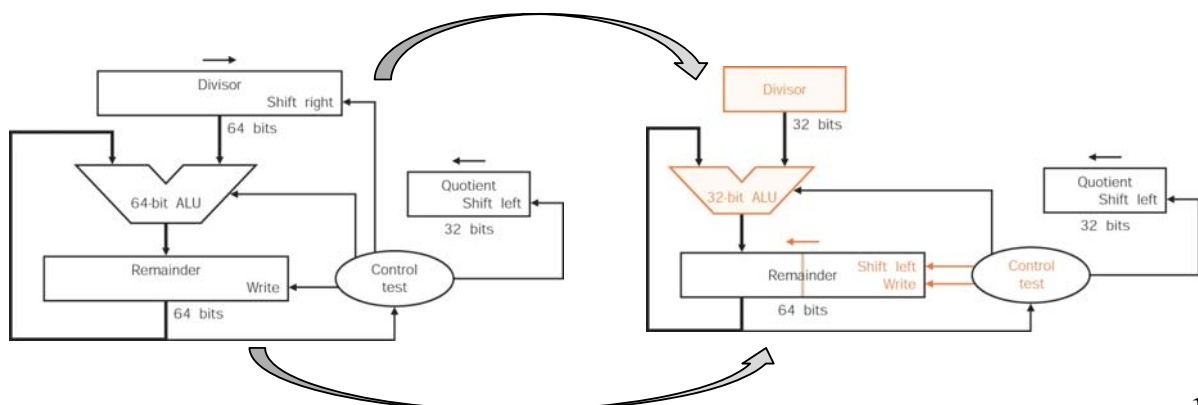
Example: 7 (0111) / 2 (0010)

Iteration	Step	Q	Divisor	Rem
0		0000	0010 0000	0000 0111
1	Rem - Div	0000	0010 0000	1110 0111
	Rest, SLQ, Q=0	0000	0010 0000	0000 0111
	Shift R Div	0000	0001 0000	0000 0111
2	Rem - Div	0000	0001 0000	1111 0111
	Rest, SLQ, Q=0	0000	0001 0000	0000 0111
	Shift R Div	0000	0000 1000	0000 0111
3	Rem - Div	0000	0000 1000	1111 1111
	Rest, SLQ, Q=0	0000	0000 1000	0000 0111
	Shift R Div	0000	0000 0100	0000 0111
4	Rem - Div	0000	0000 0100	0000 0011
	SLQ, Q=1	0001	0000 0100	0000 0011
	Shift R Div	0001	0000 0010	0000 0011
5	Rem - Div	0001	0000 0010	0000 0001
	SLQ, Q=1	0011	0000 0010	0000 0001
	Shift R Div	0011	0000 0001	0000 0001

CE, KWU Prof. S.W. LEE 15

Observations on Divide Version 1

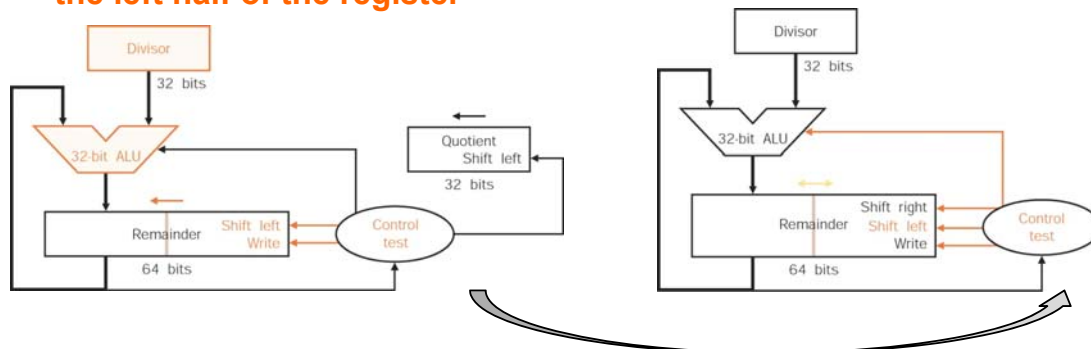
- 1/2 bits in divisor always 0
=> 1/2 of 64-bit adder is wasted
=> 1/2 of divisor is wasted
- Instead of shifting divisor to right, shift remainder to left?
- Notice that there is no way to get a 1 in leading digit! (this would be an overflow, since quotient would have n+1 bits)
 - 1st step cannot produce a 1 in quotient bit (otherwise too big)
 - switch order to shift first and then subtract, can save 1 iteration



CE, KWU Prof. S.W. LEE 16

Observations on Divide Version 2

- Eliminate Quotient register by combining with Remainder as shifted left
 - Start by shifting the Remainder left as before.
 - Thereafter loop contains only two steps because the shifting of the Remainder register shifts both the remainder in the left half and the quotient in the right half
 - The consequence of combining the two registers together and the new order of the operations in the loop is that the remainder will shifted left one time too many.
 - Thus the final correction step must **shift back only the remainder in the left half of the register**



CE, KWU Prof. S.W. LEE

17

Observations on Divide Version 3

- Same Hardware as Multiply: just need ALU to add or subtract, and 63-bit register to shift left or shift right
- Hi and Lo registers in MIPS combine to act as 64-bit register for multiply and divide
- Signed Divides: Simplest is to remember signs, make positive, and complement quotient and remainder if necessary
 - Note: Dividend and Remainder must have same sign
 - Note: Quotient negated if Divisor sign & Dividend sign disagree
e.g., $-7 \div 2 = -3$, remainder = -1
- Possible for quotient to be too large: if divide 64-bit integer by 1, quotient is 64 bit

CE, KWU Prof. S.W. LEE

18

Faster Division

- Can't use parallel hardware as in multiplier
 - Subtraction is conditional on sign of remainder
- Faster dividers (e.g. SRT division) generate multiple quotient bits per step
 - Still require multiple steps
 - SRT division: Sweeney, Robertson, and Tocher.
Lookup table based on the dividend and the divisor to determine each quotient digit.


MIPS Division

- Use HI/LO registers for result
 - HI: 32-bit remainder
 - LO: 32-bit quotient
- Instructions
 - `div rs, rt` / `divu rs, rt`
 - No overflow or divide-by-0 checking
 - Software must perform checks if required
 - Use `mfhi` , `mflo` to access result


Shifters

- Two kinds:

- Logical -- value shifted in is always "0"




- Arithmetic -- on right shifts, sign extend




- Note: these are single bit shifts. A given instruction might request 0 to 32 bits to be shifted!

Combinational Shifter from MUXes (Barrel Shifter)



- What comes in the MSBs?
- How many levels for 32-bit shifter?
- What if we use 4-1 Muxes ?

Recall Scientific Notation




Floating Point (a brief look)

- We need a way to represent
 - numbers with fractions, e.g., 3.141592
 - very small numbers, e.g., 0.000000001
 - very large numbers, e.g., 3.15576×10^9
- Representation:
 - sign, exponent, significand: $(-1)^{\text{sign}} \times \text{significand} \times 2^{\text{exponent}}$
 - more bits for **significand** gives more **accuracy**
 - more bits for **exponent** increases **range**
- IEEE 754 floating point standard:
 - single precision (32 bits): 8 bit exponent, 23 bit significand
 - double precision (64 bits): 11 bit exponent, 52 bit significand

s	exponent	significand
---	----------	-------------

IEEE 754 floating-point standard

- Leading “1” bit of significand is implicit



- Exponent is “biased” to make **sorting** easier
 - all 0s is smallest exponent all 1s is largest
 - bias of 127 for single precision and 1023 for double precision
 - summary: $(-1)^{\text{sign}} \times (1 + \text{significand}) \times 2^{\text{exponent} - \text{bias}}$

- Example:

$$\text{IEEE F.P.} \quad \pm 1.M \times 2^{e - 127}$$

- decimal: $-0.75 = -3/4 = -3/2^2 = -3 \times 2^{-2}$
- binary: $-0.11 = -1.1 \times 2^{-1}$
- floating point: exponent = 126 = 01111110
- IEEE single precision: 10111111010000000000000000000000

CE, KWU Prof. S.W. LEE 25

Converting from Binary to Decimal Floating Point

- What is the decimal single-precision floating point number that corresponds to the bit pattern **01000100010010010000000000000000**?

- Use the equation

$$X = (-1)^S \times 2^{E-127} \times (1.M)$$

where

$$S = 0$$

$$E = 10001000_2 = 136_{10}$$

$$1.M = 1.100100100000000000000000 = 1 + 2^{-1} + 2^{-4} + 2^{-7} = 1.5703125$$

so

$$X = (-1)^0 \times 2^{136-127} \times 1.5703125 = 804 = 8.04 \times 10^2$$

Converting from Decimal to Binary Floating Point

- [illegible]

CE, KWU Prof. S.W. LEE 27

Special Numbers

- **E = 0 is special for this reason.**
 - **Zero significand: number is 0**
$$(-1)^S(0.0) \times 2^{-126} = \pm 0$$
 - **Non-zero significand: denormalized number**
$$(-1)^S(0.M) \times 2^{-126}$$
 - **Denormal numbers are used to extend the range of floating-point numbers.**
- **Infinity and NaNs**
 - **Representation: exponent is all 1's (255 or 2047).**
 - **If significand is 0, ∞ ; otherwise NaN (not a number).**

Exponent	Zero Significand	Non-zero Significand	Equation
00 _H	<u>zero</u> , <u>-0</u>	<u>Subnormal Numbers</u>	$(-1)^{\text{sign}} \times 2^{-126} \times 0.\text{significand}$
01 _H , ..., FE _H	normalized value		$(-1)^{\text{sign}} \times 2^{\text{exponent}-127} \times 1.\text{significand}$
FF _H	<u>±infinity</u>	<u>NaN</u>	

CE, KWU Prof. S.W. LEE 28

Floating Point Complexities

- Operations are somewhat more complicated (see text)
- In addition to overflow we can have “underflow”
- Accuracy can be a big problem
 - IEEE 754 keeps extra bits, **guard**, **round**, and **sticky**
 - Guard bit for aligned exponents
 - Four rounding modes for round-off
 - positive divided by zero yields “infinity”
 - zero divide by zero yields “not a number”
 - other complexities
- Implementing the standard can be tricky
- Not using the standard can be even worse
 - see text for description of 80x86 and Pentium bug!

CE, KWU Prof. S.W. LEE 29

Additional Bits

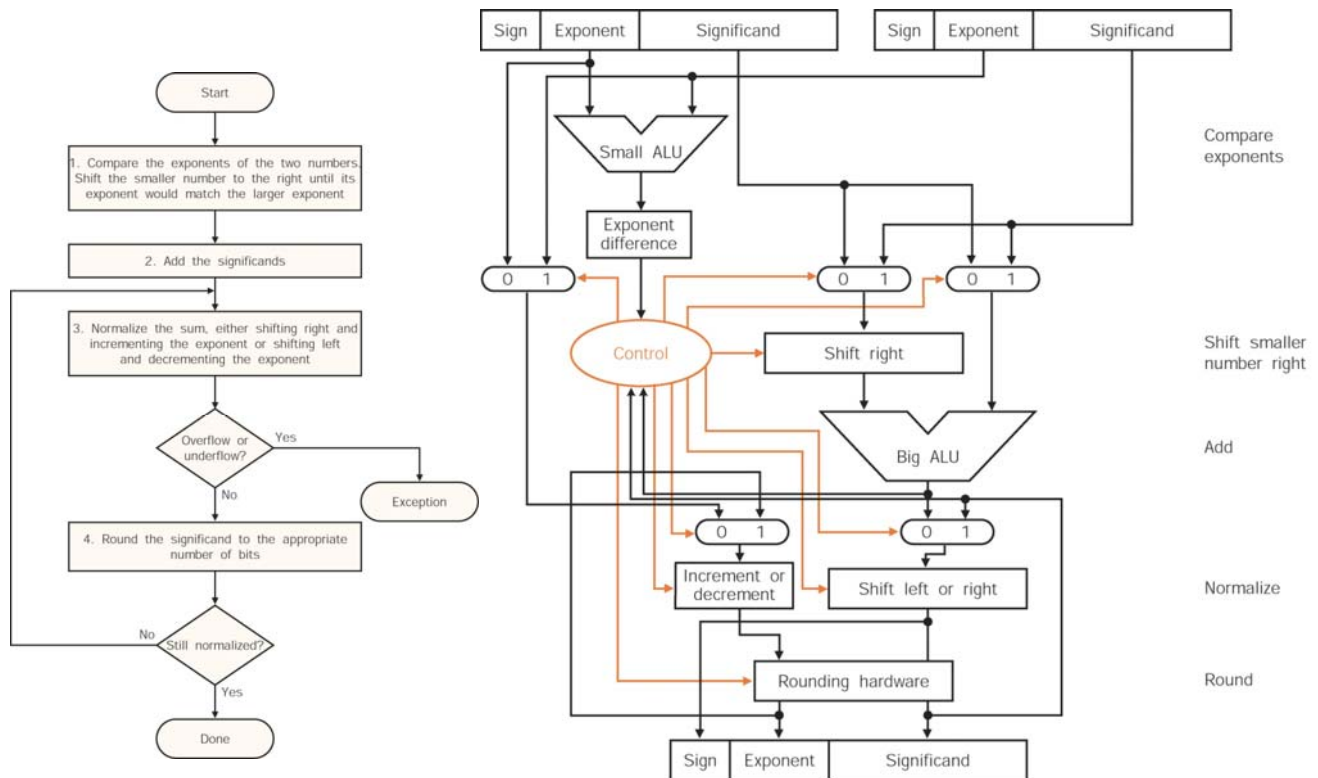
- The maximum permissible error is $\frac{1}{2}$ ulp
- Guard bit and Round bit
 - Example: 4-bit significand
$$1.0000 \times 2^0 - 1.0001 \times 2^{-2}$$
 - Align
$$\begin{array}{rcl} 1.0000 & \times & 2^0 \\ - 0.0100\textcolor{blue}{0}\textcolor{red}{1} & \times & 2^0 \end{array}$$
 - Subtract
$$0.1011\textcolor{red}{1}\textcolor{blue}{1} \times 2^0$$
 - Normalize/Round
$$\begin{array}{rcl} 1.011\textcolor{red}{1}\textcolor{blue}{1} & \times & 2^{-1} \\ 1.1000 & \times & 2^{-1} \text{ (simple round up)} \end{array}$$
 - Without extra bit, result would be 1.0111×2^{-1} .
- Sticky bit
 - keep track of whether the bits “shifted out” are non-zero.

4 rounding modes:

- round to nearest even (default)
- round toward $+\infty$
- round toward $-\infty$
- round toward 0

CE, KWU Prof. S.W. LEE 30

Addition Block Diagram



CE, KWU Prof. S.W. LEE 31

Multiplication

Consider a 4-digit binary example

$$1.000_2 \times 2^{-1} \times -1.110_2 \times 2^{-2} (0.5 \times -0.4375)$$

1. Add exponents

$$\text{Unbiased: } -1 + -2 = -3$$

$$\text{Biased: } (-1 + 127) + (-2 + 127) - 127 = -3 + 127$$

2. Multiply significands

$$1.000_2 \times 1.110_2 = 1.1102 \Rightarrow 1.110_2 \times 2^{-3}$$

3. Normalize result & check for over/underflow

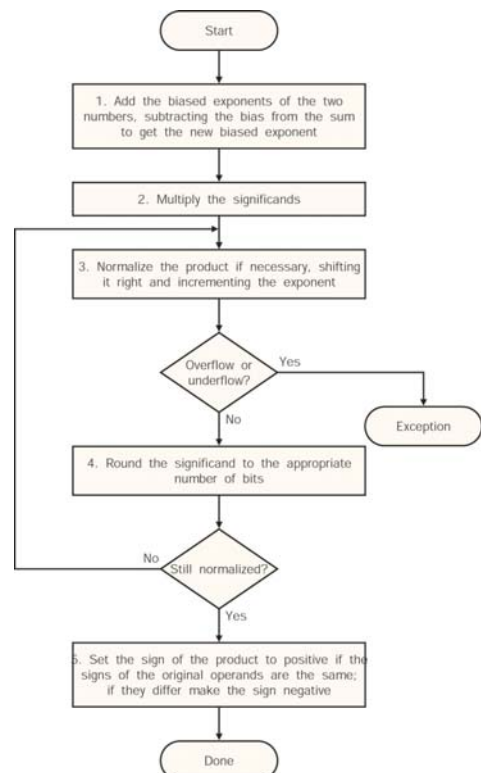
$$1.110_2 \times 2^{-3} \text{ (no change) with no over/underflow}$$

4. Round and renormalize if necessary

$$1.110_2 \times 2^{-3} \text{ (no change)}$$

5. Determine sign: +ve \times -ve \Rightarrow -ve

$$-1.110_2 \times 2^{-3} = -0.21875$$



CE, KWU Prof. S.W. LEE 32

Summary

- **Computer arithmetic is constrained by limited precision**
- **Bit patterns have no inherent meaning but standards do exist**
 - **two's complement**
 - **IEEE 754 floating point**
- **Computer instructions determine “meaning” of the bit patterns**
- **Performance and accuracy are important so there are many complexities in real machines (i.e., algorithms and implementation).**