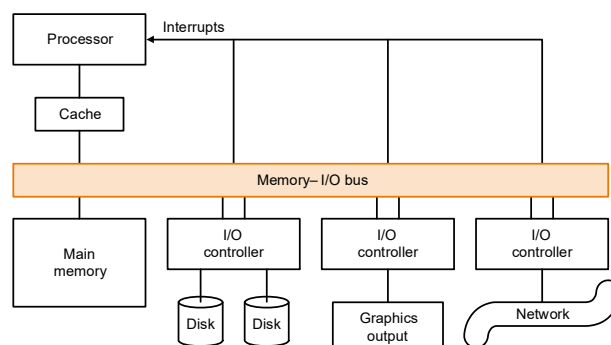

Peripherals (1/2)

1

Interfacing Processors and Peripherals

- I/O Design affected by many factors (expandability, resilience)
- Performance:
 - access latency
 - throughput
 - connection between devices and the system
 - the memory hierarchy
 - the operating system
- A variety of different users (e.g., banks, supercomputers, engineers)

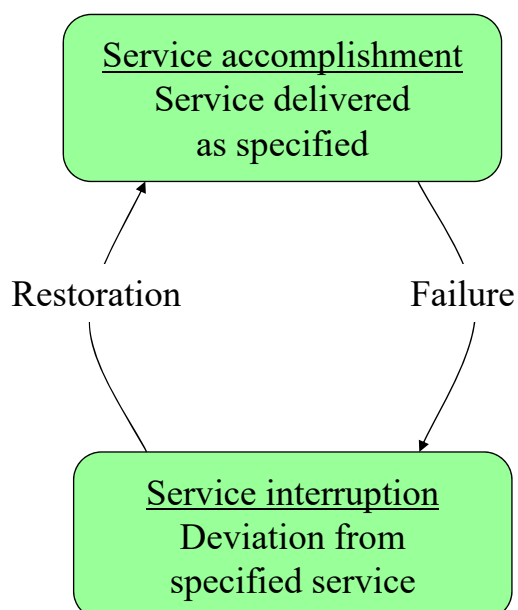


I/O System Characteristics

- **Dependability is important**
 - Particularly for storage devices
- **Performance measures**
 - Latency (response time)
 - Throughput (bandwidth)
 - **Desktops & embedded systems**
 - Mainly interested in response time & diversity of devices
 - **Servers**
 - Mainly interested in throughput & expandability of devices


Dependability

- **Fault: failure of a component**
 - May or may not lead to system failure



Dependability Measures

- Reliability: mean time to failure (MTTF)
- Service interruption: mean time to repair (MTTR)
- Mean time between failures
 - $MTBF = MTTF + MTTR$
- Availability = $MTTF / (MTTF + MTTR) = MTTF / MTBF$
- Improving Availability
 - Increase MTTF: fault avoidance, fault tolerance, fault forecasting
 - Reduce MTTR: improved tools and processes for diagnosis and repair




$$\text{Time Between Failures} = \{ \text{down time} - \text{up time} \}$$

$$\text{Mean time between failures} = MTBF = \frac{\Sigma(\text{downtime} - \text{uptime})}{\text{number of failures}}$$

CE, KWU Prof. S.W. LEE 5

Disk Storage

- Nonvolatile, rotating magnetic storage



Disk Sectors and Access

- Each sector records
 - Sector ID
 - Data (512 bytes, 4096 bytes proposed)
 - Error correcting code (ECC)
 - Used to hide defects and recording errors
 - Synchronization fields and gaps
- Access to a sector involves
 - Queuing delay if other accesses are pending
 - Seek: move the heads
 - Rotational latency
 - Data transfer
 - Controller overhead

Disk Access Example

- Given
 - 512B sector, 15,000rpm, 4ms average seek time, 100MB/s transfer rate, 0.2ms controller overhead, idle disk
- Average read time
 - 4ms seek time
 - + $\frac{1}{2} / (15,000/60) = 2\text{ms}$ rotational latency
 - + $512 / 100\text{MB/s} = 0.005\text{ms}$ transfer time
 - + 0.2ms controller delay
 - = 6.2ms
- If actual average seek time is 1ms
 - Average read time = 3.2ms

Disk Performance Issues

- **Manufacturers quote average seek time**
 - Based on all possible seeks
 - Locality and OS scheduling lead to smaller actual average seek times
- **Smart disk controller allocate physical sectors on disk**
 - Present logical sector interface to host
 - SCSI (Small Computer System Interface), ATA (Advanced Technology Attachment), SATA (Serial ATA)
- **Disk drives include caches**
 - Prefetch sectors in anticipation of access
 - Avoid seek and rotational delay

Flash Storage

- **Nonvolatile semiconductor storage**
 - $100\times - 1000\times$ faster than disk
 - Smaller, lower power, more robust
 - But more \$/GB (between disk and DRAM)



Flash Types


- **NOR flash: bit cell like a NOR gate**
 - Random read/write access
 - Used for instruction memory in embedded systems
- **NAND flash: bit cell like a NAND gate**
 - Denser (bits/area), but block-at-a-time access
 - Cheaper per GB
 - Used for USB keys, media storage, ...
- **Flash bits wears out after 1000's of accesses**
 - Not suitable for direct RAM or disk replacement
 - Wear leveling: remap data to less used blocks

RAID

- **Redundant Array of Inexpensive (Independent) Disks**
 - Use multiple smaller disks (c.f. one large disk)
 - Parallelism improves performance
 - Plus extra disk(s) for redundant data storage
- **Provides fault tolerant storage system**
 - Especially if failed disks can be “hot swapped”
- **It's a technology that enables greater levels of performance, reliability and/or large volumes when dealing with data.**
- **Mirroring, Stripping (of data) and Error correction techniques combined with multiple disk arrays give you the reliability and performance.**

RAID 0

- It splits data among two or more disks.
- Provides good performance.
- Lack of data redundancy means there is no fail over support with this configuration.
- In the diagram to the right, the odd blocks are written to disk 0 and the even blocks to disk 1 such that A1, A2, A3, A4, ... would be the order of blocks read if read sequentially from the beginning.
- Used in read only NFS systems and gaming systems.
- JBOD = Just a bunch of disks




RAID 0 analysis

- **Failure Rate:**
 - MTBF of RAID0 is roughly proportional to the number of disks in the array.
 - $\text{Pr}(\text{disk.fail}) = 5\%$, then
 $\text{Pr}(\text{at.least.one.fails}) = 1 - \text{Pr}(\text{none.fails}) = 1 - (1 - 0.05)^2 = 9.75\%$
- **Performance:**
 - The fragments are written to their respective disks simultaneously on the same sector.
 - This allows smaller sections of the entire chunk of data to be read off the drive in parallel, hence good performance.

RAID 1

- RAID1 is 'data mirroring'.
- Two copies of the data are held on two physical disks, and the data is always identical.
- Twice as many disks are required to store the same data when compared to RAID 0.
- Array continues to operate so long as at least one drive is functioning.
- Failure Rate:
 - If $\text{Pr}(\text{disk.fail}) = 5\%$, then the probability of both the drives failing in a 2-disk array is $\text{Pr}(\text{both.fail}) = (0.05)^2 = 0.25\%$.
- Performance:
 - If we use independent disk controllers for each disk, then we can increase the read or write speeds by doing operations in parallel.



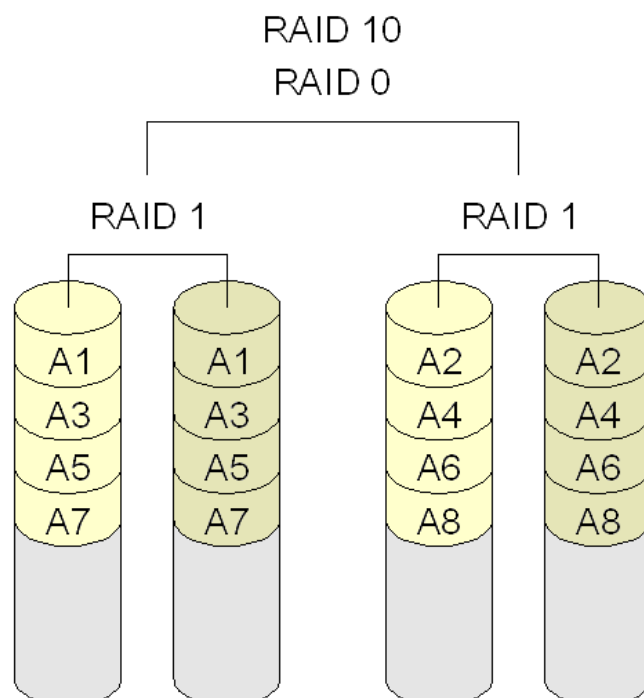
CE, KWU Prof. S.W. LEE

15

RAID 10

- Combines RAID 1 and RAID 0.
- Which means having the pleasure of both - good performance and good failover handling.
- Also called 'Nested RAID'.

- RAID 01
 - Mirroring of striped disks
 - 2nd failure while repairing results in the system failure.



CE, KWU Prof. S.W. LEE

16

RAID 2: Bit-Interleaved with ECC

- **Error correcting code (ECC)**
 - **N + E disks (e.g., 10 + 4)**
 - **Split data at bit level across N disks**
 - **Generate E-bit ECC**
 - **Too complex, not used in practice**

RAID 3: Byte-Interleaved Parity



- **N + 1 disks**
 - **Data striped across N disks at byte level**
 - **Redundant disk stores parity**
 - **Read access**
 - Read all disks
 - **Write access**
 - Generate new parity and update all disks
 - **On failure**
 - Use parity to reconstruct missing data
- **Not widely used**

RAID 4: Block-Interleaved Parity

- **N + 1 disks**
 - **Data striped across N disks at block level**
 - **Redundant disk stores parity for a group of blocks**
 - **Read access**
 - Read only the disk holding the required block
 - **Write access**
 - Just read disk containing modified block, and parity disk
 - Calculate new parity, update data disk and parity disk
 - **On failure**
 - Use parity to reconstruct missing data
- **Not widely used**

CE, KWU Prof. S.W. LEE 19


RAID 3 vs RAID 4



CE, KWU Prof. S.W. LEE 20


RAID 5: Distributed Block Parity

- **N + 1 disks**
 - Like RAID 4, but parity blocks distributed across disks
 - Avoids parity disk being a bottleneck
- **Widely used**



RAID 6: P + Q Redundancy

- **N + 2 disks**
 - Like RAID 5, but two lots of parity
 - Greater fault tolerance through more redundancy
- It is seen as the best way to guarantee data integrity as it uses double parity.
- Lesser MTBF compared to RAID5.
- It has a drawback though of longer write time.




Implementations

- **Software based RAID:**
 - Software implementations are provided by many Operating Systems.
 - A software layer sits above the disk device drivers and provides an abstraction layer between the logical drives(RAIDs) and physical drives.
 - Server's processor is used to run the RAID software.
 - Used for simpler configurations like RAID0 and RAID1.
- **Hardware based RAID:**
 - A hardware implementation of RAID requires at least a special-purpose RAID controller.
 - On a desktop system this may be built into the motherboard.
 - Processor is not used for RAID calculations as a separate controller present.

A Bus is:

- shared communication link
- single set of wires used to connect multiple subsystems



- **A Bus is also a fundamental tool for composing large, complex systems**
 - systematic means of abstraction

Bus basics

- A bus consists of
 - A set of wires over which data and bus commands are transferred
 - *Communication protocols* for determining bus ownership, (which device can currently use the bus), bus commands, and responses to commands
- To use the bus, a device (cache, memory, or I/O) must first *arbitrate* for the usage of the bus in order to get exclusive usage for the required time period
- A bus operation may consist of sending
 - Bus commands (e.g., read, write)
 - Addresses (e.g., of the memory location to read or write)
 - Data (e.g., to be returned in response to a read command)

Interconnecting Components

- Need interconnections between
 - CPU, memory, I/O controllers
- Bus: shared communication channel
 - Parallel set of wires for data and synchronization of data transfer
 - Can become a bottleneck
- Performance limited by physical factors
 - Wire length, number of connections
- More recent alternative: high-speed serial connections with switches
 - Like networks

Bus Types

- **Processor-Memory buses**
 - Short, high speed
 - Design is matched to memory organization
- **I/O buses**
 - Longer, allowing multiple connections
 - Specified by standards for interoperability
 - Connect to processor-memory bus through a bridge

Bus Signals and Synchronization

- **Data lines**
 - Carry address and data
 - Multiplexed or separate
- **Control lines**
 - Indicate data type, synchronize transactions
- **Synchronous**
 - Uses a bus clock
- **Asynchronous**
 - Uses request/acknowledge control lines for handshaking

Advantages and Disadvantage of Buses

- **Advantages**
 - **Versatility:**
 - New devices can be added easily
 - Peripherals can be moved between computer systems that use the same bus standard
 - **Low Cost:**
 - A single set of wires is shared in multiple ways
 - Manage complexity by partitioning the design
- **Disadvantages**
 - **It creates a communication bottleneck**
 - The bandwidth of that bus can limit the maximum I/O throughput
 - **The maximum bus speed is largely limited by:**
 - The **length** of the bus
 - The **number** of devices on the bus
 - The need to support a range of devices with:
 - Widely varying latencies
 - Widely varying data transfer rates

CE, KWU Prof. S.W. LEE 29

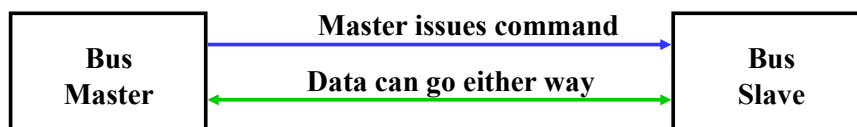
The General Organization of a Bus

- **Control lines:**
 - **Signal requests and acknowledgments**
 - **Indicate what type of information is on the data lines**
- **Data lines:**
 - **carry information between the source and the destination:**
 - Data and Addresses
 - Complex commands
- **What defines a bus?**
 - **Transaction Protocol**
 - **Timing and Signaling Specification**
 - **Bunch of Wires**
 - **Electrical Specification**
 - **Physical / Mechanical Characteristics – the connectors**

CE, KWU Prof. S.W. LEE 30

Master versus Slave


- A bus transaction includes two parts:
 - Issuing the command (and address) – request
 - Transferring the data – action
- Master is the one who starts the bus transaction by:
 - issuing the command (and address)
- Slave is the one who responds to the address by:
 - Sending data to the master if the master ask for data
 - Receiving data from the master if the master wants to send data



Types of Buses


- **Processor-Memory Bus (design specific)**
 - Short and high speed
 - Only need to match the memory system
 - Maximize memory-to-processor bandwidth
 - Connects directly to the processor
 - Optimized for cache block transfers
- **I/O Bus (industry standard)**
 - Usually is lengthy and slower
 - Need to match a wide range of I/O devices
 - Connects to the processor-memory bus or backplane bus
- **Backplane Bus (standard or proprietary)**
 - Backplane: an interconnection structure within the chassis
 - Allow processors, memory, and I/O devices to coexist
 - Cost advantage: one bus for all components

A Computer System with One Bus: Backplane Bus




- A single bus (the backplane bus) is used for:
 - Processor to memory communication
 - Communication between I/O devices and memory
- Advantages: Simple and low cost
- Disadvantages: slow and the bus can become a major bottleneck
- Example: IBM PC - AT

A Two-Bus System



- I/O buses tap into the processor-memory bus via bus adaptors:
 - Processor-memory bus: mainly for processor-memory traffic
 - I/O buses: provide expansion slots for I/O devices
- Apple Macintosh-II
 - NuBus: Processor, memory, and a few selected I/O devices
 - SCSI Bus: the rest of the I/O devices

A Three-Bus System



- A small number of backplane buses tap into the processor-memory bus
 - Processor-memory bus is used for processor memory traffic
 - I/O buses are connected to the backplane bus
- Advantage: loading on the processor bus is greatly reduced