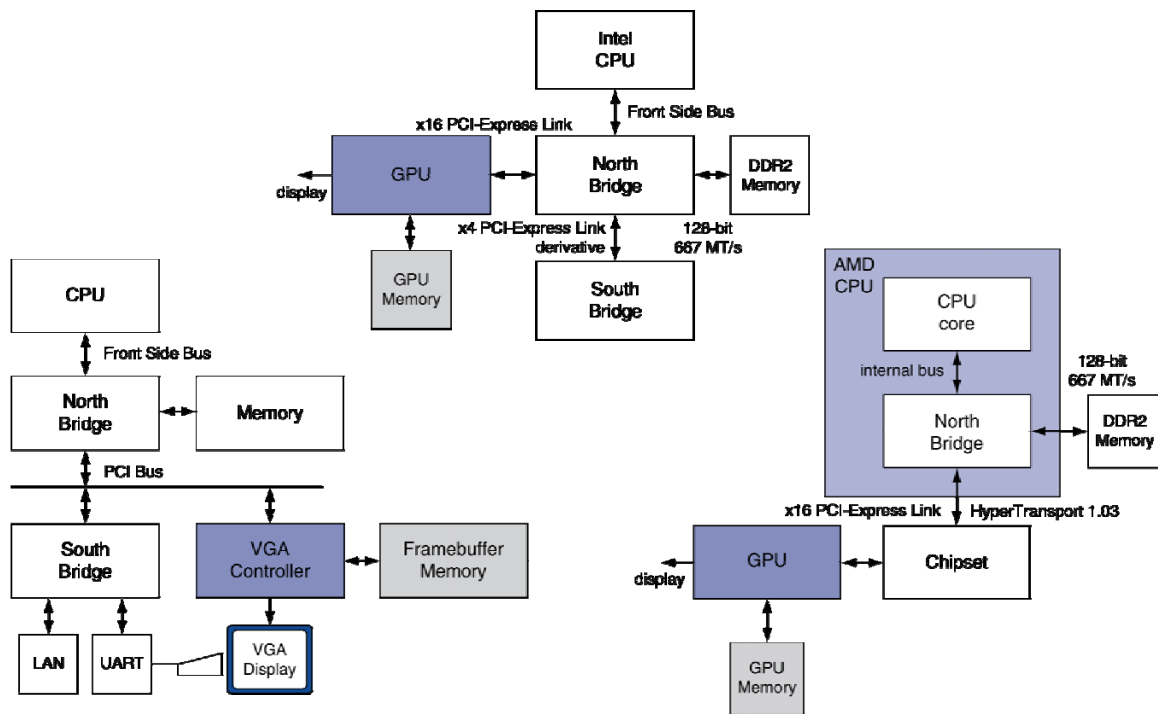

Chapter Six (2/2)

1

History of GPUs

- **Early video cards**
 - Frame buffer memory with address generation for video output
- **3D graphics processing**
 - Originally high-end computers (e.g., SGI)
 - Moore's Law \Rightarrow lower cost, higher density
 - 3D graphics cards for PCs and game consoles
- **Graphics Processing Units**
 - Processors oriented to 3D graphics tasks
 - Vertex/pixel processing, shading, texture mapping, rasterization

Graphics in the System



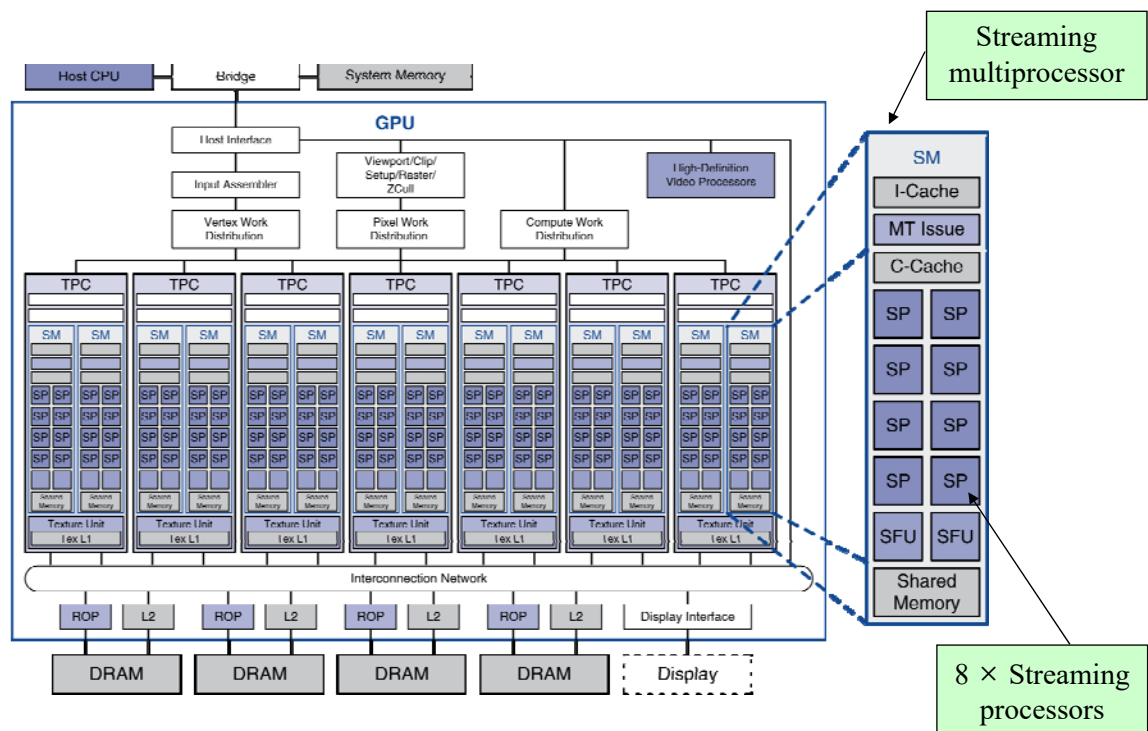
CE, KWU Prof. S.W. LEE 3

GPU Architectures

- **Processing is highly data-parallel**
 - GPUs are highly multithreaded
 - Use thread switching to hide memory latency
 - Less reliance on multi-level caches
 - Graphics memory is wide and high-bandwidth
- **Trend toward general purpose GPUs**
 - Heterogeneous CPU/GPU systems
 - CPU for sequential code, GPU for parallel code
- **Programming languages/APIs**
 - DirectX, OpenGL
 - C for Graphics (Cg), High Level Shader Language (HLSL)
 - Compute Unified Device Architecture (CUDA)
 - OpenCL

CE, KWU Prof. S.W. LEE 4

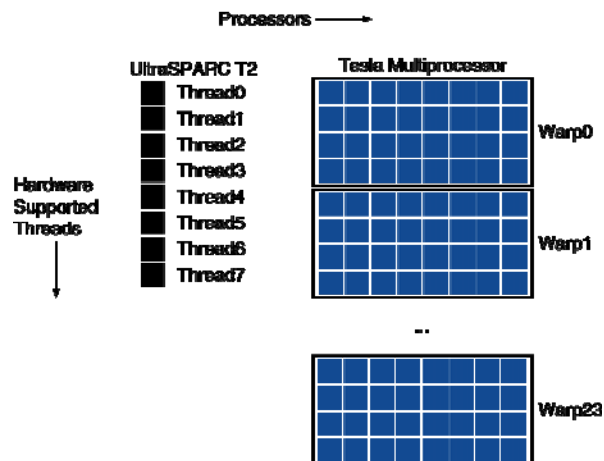
Example: NVIDIA Tesla



CE, KWU Prof. S.W. LEE 5

Example: NVIDIA Tesla

- Streaming Processors
 - Single-precision FP and integer units
 - Each SP is fine-grained multithreaded
- Warp: group of 32 threads
 - Executed in parallel, SIMD style
 - 8 SPs
 - × 4 clock cycles
 - Hardware contexts for 24 warps
 - Registers, PCs, ...



CE, KWU Prof. S.W. LEE 6

Classifying GPUs

- **Don't fit nicely into SIMD/MIMD model**
 - **Conditional execution in a thread allows an illusion of MIMD**
 - But with performance degradation
 - Need to write general purpose code with care

	Static: Discovered at Compile Time	Dynamic: Discovered at Runtime
Instruction-Level Parallelism	VLIW	Superscalar
Data-Level Parallelism	SIMD or Vector	Tesla Multiprocessor

Coherence Defined

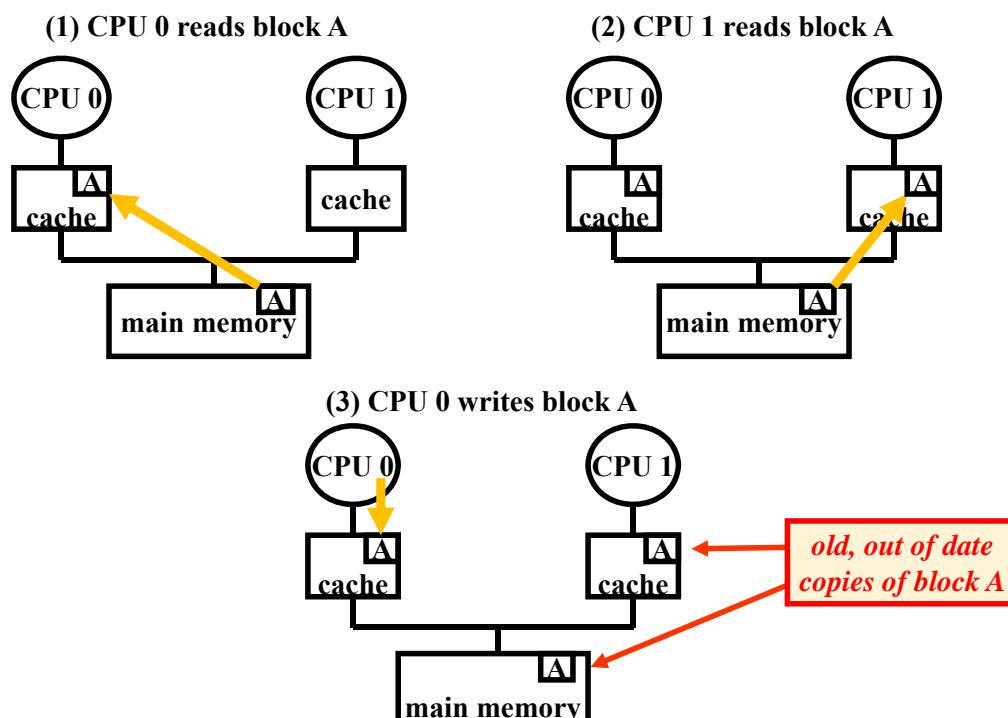
- **Informally: Reads return most recently written value**
- **Formally:**
 - **P writes X; P reads X (no intervening writes)**
⇒ **read returns written value**
 - **P1 writes X; P2 reads X (sufficiently later)**
⇒ **read returns written value**
 - c.f. CPU B reading X after step 3 in example
 - **P1 writes X, P2 writes X**
⇒ **all processors see writes in the same order**
 - End up with the same final value for X

Cache Coherence Protocols

- Operations performed by caches in multiprocessors to ensure coherence
 - Migration of data to local caches
 - Reduces bandwidth for shared memory
 - Replication of read-shared data
 - Reduces contention for access
- Snooping protocols
 - Each cache monitors bus reads/writes
- Directory-based protocols
 - Caches and memory record sharing status of blocks in a directory

Cache Coherence Problem

- Two writeback caches becoming incoherent



Cache Coherence Problem

- Suppose two CPU cores share a physical address space
 - Write-through caches

Time step	Event	CPU A's cache	CPU B's cache	Memory X
0				0
1	CPU A reads X	0		0
2	CPU B reads X	0	0	0
3	CPU A writes 1 to X	1	0	1

Cache coherence protocols

- Ensures that cached blocks that are written to are observable by all processors
- Assigns a *state* field to all cached blocks
- Defines actions for performing reads and writes to blocks in each state that ensure cache coherence
- Actions are much more complicated than described here in a real machine with a split transaction bus

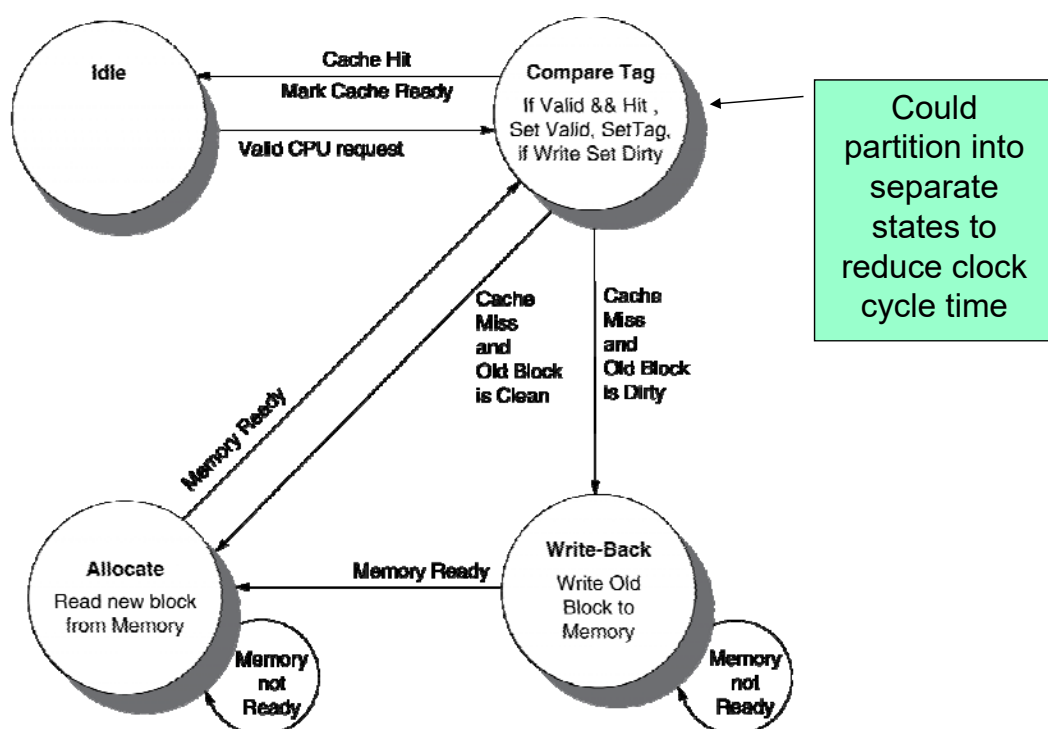
Invalidating Snooping Protocols

- Cache gets exclusive access to a block when it is to be written
 - Broadcasts an invalidate message on the bus
 - Subsequent read in another cache misses
 - Owning cache supplies updated value

CPU activity	Bus activity	CPU A's cache	CPU B's cache	Memory X
				0
CPU A reads X	Cache miss for X	0		0
CPU B reads X	Cache miss for X	0	0	0
CPU A writes 1 to X	Invalidate for X	1		0
CPU B read X	Cache miss for X	1	1	1

CE, KWU Prof. S.W. LEE 13

Cache Controller FSM for a Single Processor



MESI cache coherence protocol

- Commonly used (or variant thereof) in shared memory multiprocessors
- Idea is to ensure that when a cache wants to write to a cache block that other remote caches *invalidate* their copies first
- Each cache block is in one of four states (2 bits stored with each cache block)
 - **I**nvalid: contents are not valid
 - **S**hared: other processor caches *may* have the same copy; main memory has the same copy
 - **E**xclusive: no other processor cache has a copy; main memory has the same copy
 - **M**odified: no other processor cache has a copy; main memory has an old copy

MESI cache coherence protocol

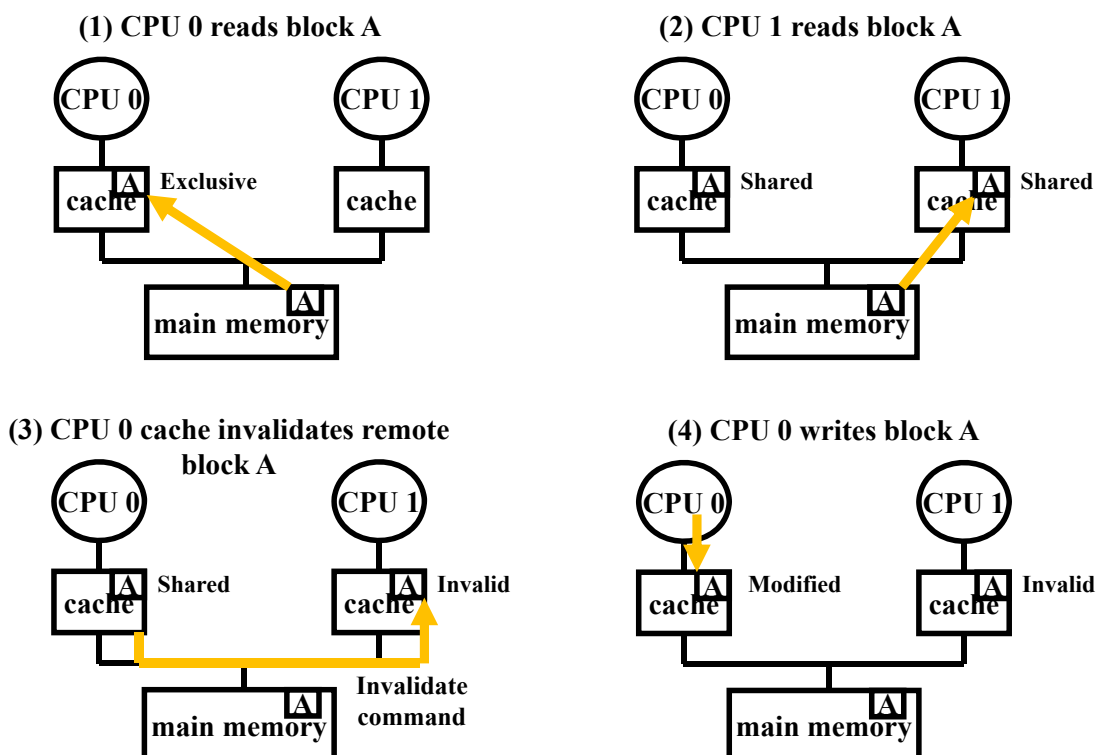
- **Cache read actions**
 - **Hit – local cache actions**
 - Read block
 - **Hit – remote cache actions**
 - None
 - **Miss – local cache actions**
 - Request block from bus
 - If not in a remote cache, set state to **Exclusive**
 - If also in a remote cache, set state to **Shared**
 - **Miss – remote cache actions**
 - Look up cache tags to see if the block is present
 - If so, signal the local cache that we have a copy, provide it if it is in state *Modified*, and change the state of our copy to **Shared**

MESI cache coherence protocol

- **Cache write actions**
 - **Hit – local cache actions**
 - Check state of block
 - If Shared, send an **Invalidation** bus command to all remote caches
 - Write the block and change the state to **Modified**
 - **Hit – remote cache actions**
 - Upon receipt of an Invalidation command on the bus, look up cache tags to see if the block is present
 - If so, change the state of the block to **Invalid**
 - **Miss – local cache actions**
 - Simultaneously **request** block from bus and send an **Invalidation** command
 - After block received, write the block and set the state to **Modified**
 - **Miss – remote cache actions**
 - Look up cache tags to see if the block is present
 - If so, signal the local cache that we have a copy, provide it if it is in state **Modified**, and change the state of our copy to **Invalid**

CE, KWU Prof. S.W. LEE 17

Cache coherence problem revisited



CE, KWU Prof. S.W. LEE 18

Memory Consistency

- When are writes seen by other processors
 - “Seen” means a read returns the written value
 - Can’t be instantaneously
- Assumptions
 - A write completes only when all processors have seen it
 - A processor does not reorder writes with other accesses
- Consequence
 - P writes X then writes Y
⇒ all processors that see new Y also see new X
 - Processors can reorder reads, but not writes
- Example
 - P1: A=1; B=1; while B=1 do nothing; print A;
 - P2: A=0; B=0;

Parallel Benchmarks

- Linpack: matrix linear algebra
- SPECrate: parallel run of SPEC CPU programs
 - Job-level parallelism
- SPLASH: Stanford Parallel Applications for Shared Memory
 - Mix of kernels and applications, strong scaling
- NAS (NASA Advanced Supercomputing) suite
 - computational fluid dynamics kernels
- PARSEC (Princeton Application Repository for Shared Memory Computers) suite
 - Multithreaded applications using Pthreads and OpenMP

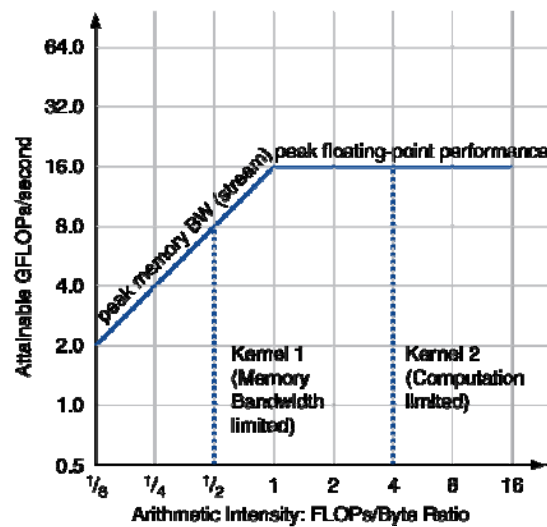
Code or Applications?

- **Traditional benchmarks**
 - Fixed code and data sets
- **Parallel programming is evolving**
 - Should algorithms, programming languages, and tools be part of the system?
 - Compare systems, provided they implement a given application
 - E.g., Linpack, Berkeley Design Patterns
- **Would foster innovation in approaches to parallelism**

Modeling Performance

- **Assume performance metric of interest is achievable GFLOPs/sec**
 - Measured using computational kernels from Berkeley Design Patterns
- **Arithmetic intensity of a kernel**
 - FLOPs per byte of memory accessed
- **For a given computer, determine**
 - Peak GFLOPS (from data sheet)
 - Peak memory bytes/sec (using Stream benchmark)


Roofline Diagram



Attainable GPLOPs/sec
 = Max (Peak Memory BW \times Arithmetic Intensity, Peak FP Performance)

Comparing Systems

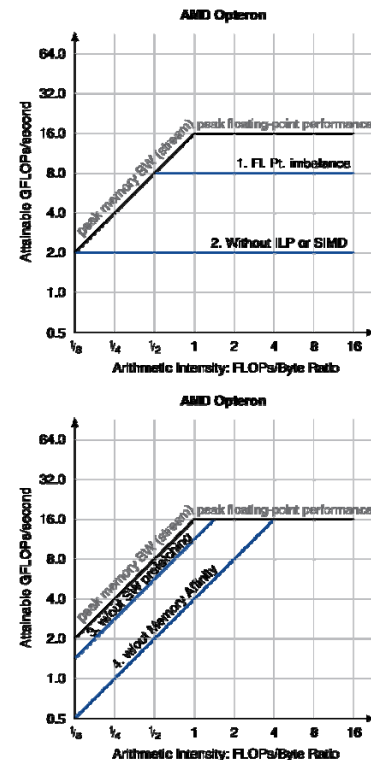
- Example: Opteron X2 vs. Opteron X4
 - 2-core vs. 4-core, $2\times$ FP performance/core, 2.2GHz vs. 2.3GHz
 - Same memory system



- To get higher performance on X4 than X2
 - Need high arithmetic intensity
 - Or working set must fit in X4's 2MB L-3 cache

Optimizing Performance

- Optimize FP performance
 - Balance adds & multiplies
 - Improve superscalar ILP and use of SIMD instructions
- Optimize memory usage
 - Software prefetch
 - Avoid load stalls
 - Memory affinity
 - Avoid non-local data accesses

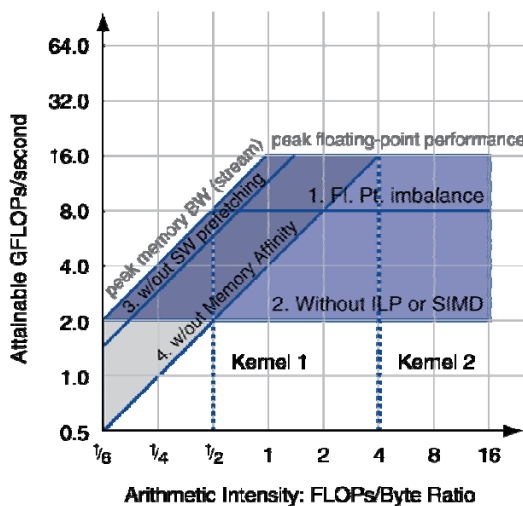


CE, KWU Prof. S.W. LEE

25

Optimizing Performance

- Choice of optimization depends on arithmetic intensity of code




- Arithmetic intensity is not always fixed
 - May scale with problem size
 - Caching reduces memory accesses
 - Increases arithmetic intensity

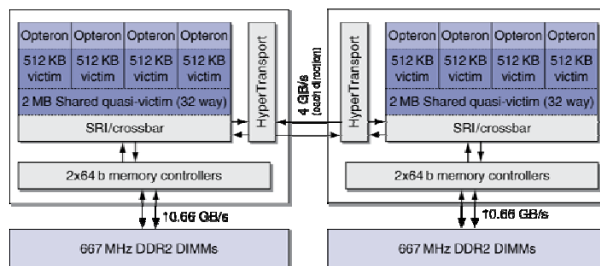
CE, KWU Prof. S.W. LEE

26

Four Example Systems



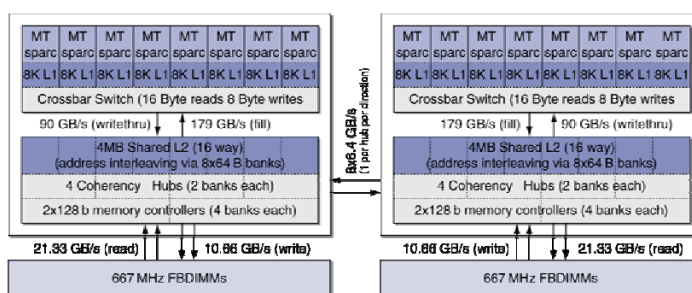
2 × quad-core
Intel Xeon e5345
(Clovertown)



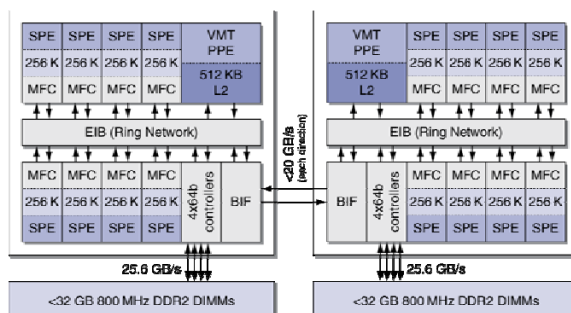
2 × quad-core
AMD Opteron X4 2356
(Barcelona)

CE, KWU Prof. S.W. LEE 27

Four Example Systems



2 × oct-core
Sun UltraSPARC
T2 5140 (Niagara 2)

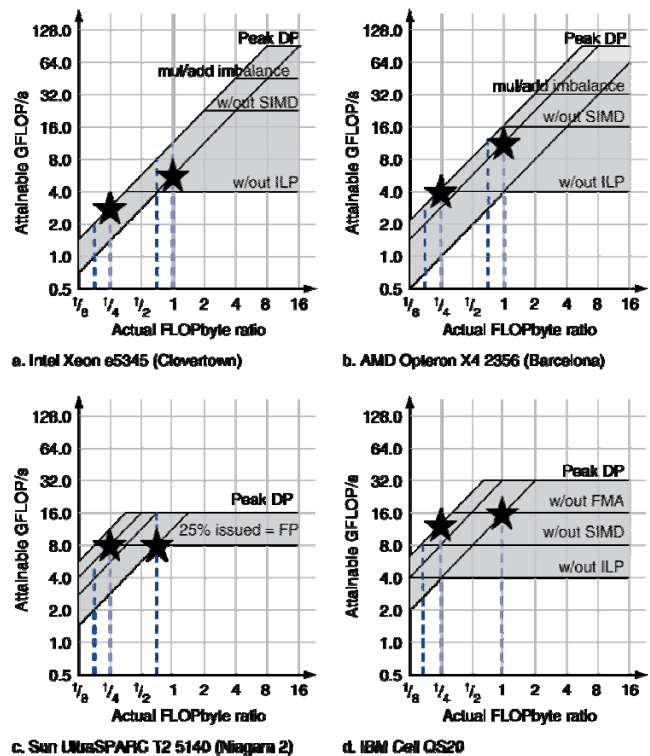


2 × oct-core
IBM Cell QS20

CE, KWU Prof. S.W. LEE 28

And Their Rooflines

- Kernels
 - SpMV (left)
 - LBHMD (right)
- Some optimizations change arithmetic intensity
- x86 systems have higher peak GFLOPs
 - But harder to achieve, given memory bandwidth

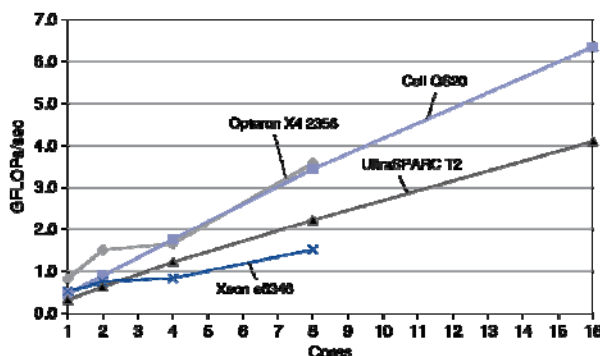


CE, KWU Prof. S.W. LEE

29

Performance on SpMV

- Sparse matrix/vector multiply
 - Irregular memory accesses, memory bound
- Arithmetic intensity
 - 0.166 before memory optimization, 0.25 after



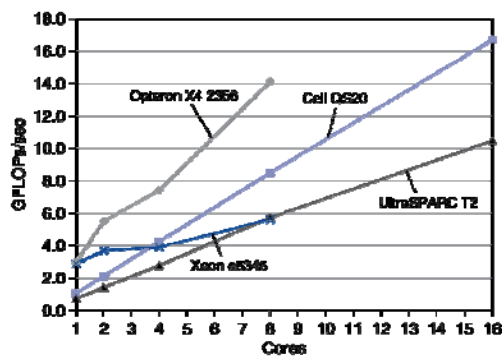
- Xeon vs. Opteron
 - Similar peak FLOPS
 - Xeon limited by shared FSBs and chipset
- UltraSPARC/Cell vs. x86
 - 20 – 30 vs. 75 peak GFLOPs
 - More cores and memory bandwidth

CE, KWU Prof. S.W. LEE

30

Performance on LBMHD

- Fluid dynamics: structured grid over time steps
 - Each point: 75 FP read/write, 1300 FP ops
- Arithmetic intensity
 - 0.70 before optimization, 1.07 after



- Opteron vs. UltraSPARC
 - More powerful cores, not limited by memory bandwidth
- Xeon vs. others
 - Still suffers from memory bottlenecks

Achieving Performance

- Compare naïve vs. optimized code
 - If naïve code performs well, it's easier to write high performance code for the system

System	Kernel	Naïve GFLOPs/sec	Optimized GFLOPs/sec	Naïve as % of optimized
Intel Xeon	SpMV	1.0	1.5	64%
	LBMHD	4.6	5.6	82%
AMD Opteron X4	SpMV	1.4	3.6	38%
	LBMHD	7.1	14.1	50%
Sun UltraSPARC T2	SpMV	3.5	4.1	86%
	LBMHD	9.7	10.5	93%
IBM Cell QS20	SpMV	Naïve code not feasible	6.4	0%
	LBMHD	Naïve code not feasible	16.7	0%

Fallacies

- **Amdahl's Law doesn't apply to parallel computers**
 - Since we can achieve linear speedup
 - But only on applications with weak scaling
- **Peak performance tracks observed performance**
 - Marketers like this approach!
 - But compare Xeon with others in example
 - Need to be aware of bottlenecks

Pitfalls

- **Not developing the software to take account of a multiprocessor architecture**
 - **Example: using a single lock for a shared composite resource**
 - Serializes accesses, even if they could be done in parallel
 - Use finer-granularity locking

Concluding Remarks

- **Goal: higher performance by using multiple processors**
- **Difficulties**
 - **Developing parallel software**
 - **Devising appropriate architectures**
- **Many reasons for optimism**
 - **Changing software and application environment**
 - **Chip-level multiprocessors with lower latency, higher bandwidth interconnect**
- **An ongoing challenge for computer architects!**